


Captioning: accuracy vs energy per caption (Colab T4, batch 1, incl. idle)



- FP32
- Decoder FP16
- Both FP16
- FP32 greedy
- Encoder FP16
- Both INT8
- FP32 beam 5
- hollow = dominated
- Encoder INT8
- Decoder NF4
- Both NF4
- Encoder NF4